Project Report

on

Non-Profit Support Triage Agent

Course Name: AgenticAI

Institution Name: Medicaps University – Datagami Skill Based Course

Submitted by:

Sr no	Student Name	Enrolment Number
1.	Ananya Subramanya Rao	EN22CS301120
2.	Anirudh Vyas	EN22CS301131
3.	Anushka Kashyap	EN22CS301180
4.	Atharv Chaturvedi	EN22CS301228
5.	Harshit Panchal	EN23CS3L1010

Group Name: Group 03D2

Project Number: AAI-15

Industry Mentor Name: Ashwin Krishnan Sir

University Mentor Name: Prof. Ajeet Singh Rajput

Academic Year: 2026

Table of Contents

Problem Statement & Objectives

1. Introduction
2. Problem Statement
3. Project Objectives
4. Scope of the Project

Proposed Solution

1. Key features
2. Overall Architecture and Workflow
3. Tools & Technologies Used

Results & Output

1. Screenshots and Outputs
2. Reports, Dashboards and Models
3. Key outcomes

Conclusion

Future Scope & Enhancements

Problem Statement & Objectives

1. Introduction

In today's fast-paced digital world, non-profit organizations receive hundreds of messages every day — donation inqueries, volunteer sign-ups, complaints, event inquiries, and partnership requests — all arriving through emails, web forms, and contact pages. Managing this volume manually is time-consuming, error-prone, and leads to delayed responses, donor dissatisfaction, overburdened staff, and poor communication with donors and volunteers.

TriageAgent is a production-grade, AI-powered communication triage system built to solve this problem. It works like a smart AI receptionist for a non-profit organization. Instead of a staff member spending 15–20 minutes reading, categorizing, drafting a response, and routing each message to the right department, TriageAgent does all of this automatically in 5–10 seconds. It leverages the power of Large Language Models and NLP to automate the initial response workflow for non-profit communications. The system processes every incoming message in real time, classifying it by urgency and intent, extracting critical data points, generating a professional draft response, and hands it to the right person for final approval.

The system is built on a modern full-stack architecture using **FastAPI** on the backend and **React 18** on the frontend. It integrates **Google Gemini 2.5 Flash** as its core Large Language Model, **spaCy** for fast offline Named Entity Recognition, **LangChain** for multi-agent orchestration, and **Firebase** for authentication and real-time cloud database storage.

At its core, TriageAgent follows a **Multi-Agent AI Architecture** where four specialized AI agents — Classifier, NER, Response Generator, and Router — work together under a central Orchestrator to process every incoming message through a defined pipeline. The system does not replace human staff — instead it assists them by doing the time-consuming work of reading, classifying, extracting key data, and drafting a response, while leaving the final approval decision to a human staff member. This **human-in-the-loop** design ensures accuracy, empathy, and reliability in every communication.

2. Problem Statement

Create a production-grade reactive triage agent for Non-Profit communications, utilizing Large Language Models to automate initial response workflows. The system performs real-time classification of incoming text to determine urgency and intent, applies Named Entity Recognition (NER) to extract critical data points such as IDs, dates, and generates contextually accurate draft responses using an LLM API to accelerate resolution times.

3. Project Objectives

The TriageAgent project was designed and built to achieve the following core objectives:

1) **Real-Time Message Classification** Automatically classify every incoming message by urgency level (Critical / High / Normal / Low) and intent category (Donation / Volunteer / Complaint / Question / Feedback / Partnership / Event / Other) using Google Gemini 2.5 Flash with a target classification accuracy of 90–95%.
2) **Named Entity Recognition (NER)** Extract critical data points from message text using a hybrid NER approach combining spaCy (for fast offline extraction of names, dates, organizations, and monetary amounts) and Gemini LLM (for context-aware extraction of domain-specific entities such as receipt IDs, donation IDs, case numbers, email addresses, and phone numbers).
3) **Automated Draft Response Generation** Generate professional, empathetic, and contextually accurate draft email responses tailored to the message content, urgency level, and extracted entities — using Gemini 2.5 Flash with a structured prompt designed to produce complete Subject line + email body output.
4) **Intelligent Department Routing** Automatically route every triaged message to the correct department (Donations, Volunteers, Support, Partnerships, Events, General) and assign a priority level (1=Critical to 4=Low) based on urgency and intent classification.
5) **Human-in-the-Loop Review** Provide a secure admin dashboard where authorized staff can review AI-generated classifications, inspect extracted entities, read draft responses, and approve or reject each response before it is sent — ensuring no AI-generated communication reaches stakeholders without human oversight.
6) **Analytics and Insights** Deliver real-time analytics on message volume, urgency distribution, intent trends, system performance metrics (average processing time, approval rate), and trending categories — enabling managers to make data-driven staffing and resource allocation decisions.
7) **Production-Grade Reliability** Build a scalable, secure system with JWT-based authentication, Firebase Auth integration, rate limiting (60 requests/minute per IP), exponential backoff retry logic for all Gemini API calls (3 attempts: 2s → 4s → 6s), Pydantic V2 input validation at every API boundary, and structured JSON logging throughout the backend.

The overall goal is to create a **production-grade, intelligent triage system** that transforms how non-profit organizations handle incoming communications — reducing manual effort, accelerating response times, improving consistency, and allowing staff to focus on high-value human interactions rather than repetitive message processing.

4. Scope of the Project

The scope of TriageAgent — Non-Profit Support System covers the following areas:

A. Functional Scope: The functional scope defines what the TriageAgent system does the features, workflows, and user-facing capabilities that are included within the boundaries of this project.

- Secure user registration and login using Firebase Authentication with JWT-based session management.
- Message input interface where staff can paste any incoming text communication along with optional sender name and email.
- Real-time AI classification of every incoming message by urgency level (Critical / High / Normal / Low) and intent category (Donation / Volunteer / Complaint / Question / Feedback / Partnership / Event / Other).
- Hybrid Named Entity Recognition (NER) to extract critical data points including names, dates, amounts, emails, phone numbers, receipt IDs, donation IDs, and case numbers.
- AI-generated professional draft email responses with subject line and complete email body tailored to message content and urgency.
- Intelligent department routing with priority assignment (1–4) based on urgency and intent classification.
- Complete triage result display across three tabs — Overview (classification + routing), Entities (extracted data pills), and Response (draft email).
- Human-in-the-loop review dashboard where staff can inspect AI results and approve or reject every draft response before sending.
- Status lifecycle management for every triage result — Pending Review → Approved / Rejected → Sent → Archived.
- Dashboard page showing all triage results with filters by urgency, status, and limit. Real-time analytics showing urgency distribution (PieChart), intent trends (BarChart), approval rate, and average processing time.

B. Technical Scope

- **Frontend:** React 18 with Vite, Tailwind CSS 3, Zustand (state management), Axios (HTTP client), Framer Motion (animations), Recharts (analytics charts).
- **Backend:** Python 3.11 with FastAPI, Uvicorn ASGI server, Pydantic V2 (request/response validation), LangChain (multi-agent orchestration).
- **AI Pipeline:** Google Gemini 2.5 Flash (LLM) for message classification, entity extraction, and draft response generation with exponential backoff retry logic.
- **NER Layer:** Hybrid approach — spaCy (offline, <50ms) for standard entities + Gemini LLM for domain-specific entities (IDs, numbers, emails).

- **Authentication:** Firebase Authentication (email/password) with dual-token model — Firebase Auth for identity management + backend JWT.
- **Database:** Firebase Firestore (Cloud NoSQL) for storing all triage results and user profiles via Firebase Admin SDK (server-side only).
- **Security:** JWT middleware, CORS policy (restricted to localhost:5173), rate limiting (60 requests/minute per IP), Pydantic V2 input validation, and environment variable-based secret management.
- **API:** 13 REST API endpoints across 4 route groups — Auth, Triage, Messages, and Analytics.

The project is currently scoped as a v1.0 web application running in a local development environment.

Out of Scope

Features such as direct email sending integration, mobile application support, multi-language input, CRM/helpdesk integration, real-time WebSocket updates, and a dedicated caching layer are not included in the current version and are planned for future development.

Proposed Solution

TriageAgent is an AI-powered communication triage system that automates the initial response workflow for non-profit organizations. Instead of staff manually reading, categorizing, and responding to every incoming message, TriageAgent processes each message through a 4-agent AI pipeline in 5–10 seconds — classifying urgency and intent, extracting critical entities, generating a professional draft response, and routing it to the correct department — all with human approval before any action is taken.

1. Key features

The TriageAgent — Non-Profit Support System includes the following key features:

- **Real-Time AI Message Classification** — Every incoming message is automatically classified by urgency level (Critical / High / Normal / Low) and intent category (Donation / Volunteer / Complaint / Question / Feedback / Partnership / Event / Other) using Google Gemini 2.5 Flash, with a confidence score and AI reasoning explanation returned for every result.
- **Hybrid Named Entity Recognition (NER)** — A two-pass hybrid NER system combines spaCy for fast offline extraction of standard entities (names, dates, organizations, amounts) with Gemini LLM for context-aware extraction of domain-specific entities (receipt IDs, donation IDs, case numbers, email addresses, phone numbers) — delivering both speed and accuracy in a single pipeline.
- **AI-Powered Draft Response Generation** — Gemini 2.5 Flash generates a complete, professional draft email response for every triaged message — including a subject line and full email body — tailored to the message content, urgency level, intent category, and extracted entities.
- **Intelligent Department Routing** — Every message is automatically routed to the correct department (Donations, Volunteers, Support, Partnerships, Events, General) and assigned a priority level from 1 (Critical) to 4 (Low) based on classification results — eliminating manual routing decisions entirely.
- **Multi-Agent AI Architecture** — The AI pipeline is built using LangChain and consists of four specialized agents — ClassifierAgent, NERAgent, ResponseAgent, and RouterAgent — coordinated by a central TriageOrchestrator. Each agent has a single responsibility and an independent fallback, so a failure in one agent never crashes the entire pipeline.
- **Human-in-the-Loop Review Dashboard** — No AI-generated response is ever sent without explicit staff approval. The admin dashboard presents every triage result in a 3-tab interface (Overview, Entities, Draft Response) where staff can inspect, approve, or reject the AI output before any communication reaches stakeholders.

- **Complete Message Management** — A filterable, paginated message dashboard allows staff to view all triaged messages, filter by urgency and status, and track the full status lifecycle — Pending Review → Approved / Rejected → Sent → Archived — with all changes timestamped in Firestore.
- **Real-Time Analytics & Insights** — The analytics dashboard provides KPI cards (7-day and 30-day message totals, approval rate, average processing time), urgency distribution PieChart, intent breakdown BarChart, and configurable time period filters — giving managers actionable data for staffing and resource decisions.
- **Secure Authentication System** — A dual-token authentication model combines Firebase Authentication for identity management and a custom JWT (HS256, 30-minute expiry) for stateless API access — with protected routes, automatic token refresh handling, and secure logout across all sessions.
- **Production-Grade Reliability** — The system includes Pydantic V2 input validation at every API boundary, rate limiting (60 requests/minute per IP), exponential backoff retry logic for all Gemini API calls (3 attempts: 2s → 4s → 6s), structured JSON logging throughout the backend, and predefined fallback responses for every agent failure scenario.
- **Animated, Responsive Frontend** — A fully responsive React 18 + Vite SPA with Framer Motion animations, a public landing page, and a complete admin dashboard accessible from any modern browser — with smooth page transitions, toast notifications, and an intuitive 3-tab triage result interface.

2. Overall Architecture and Workflow

A. Overall Architecture - TriageAgent is built on a **3-Tier Client-Server Architecture** with an embedded **Multi-Agent AI Pipeline** in the backend tier.

Layer 1 — Presentation Layer (Frontend) The frontend is a Single Page Application (SPA) built with React 18 and Vite. All application state is managed client-side using three Zustand stores — useAuthStore (persisted to localStorage), useTriageStore, and useMessageStore. The frontend communicates with the backend exclusively through Axios HTTP calls with a JWT Bearer token attached automatically by an Axios request interceptor on every outgoing request. React Router v7 enforces protected routes — any unauthenticated user is immediately redirected to the login page.

Layer 2 — Business Logic Layer (Backend) The backend is a FastAPI application running on Uvicorn ASGI server. It exposes 13 REST API endpoints across four route modules — triage.py, messages.py, analytics.py, and auth.py. All incoming requests pass through JWT middleware, CORS middleware, and a rate limiter before reaching any route handler. The central component of this layer is the TriageOrchestrator — a LangChain-powered coordinator that sequences four AI agents in a defined pipeline for every triage request.

Layer 3 — AI / Agent Service Layer The AI workload is handled by a four-agent pipeline embedded within the backend tier. Each agent is a single-responsibility class:

Agent	Responsibility
ClassifierAgent	Classifies urgency (Critical/High/Normal/Low) and intent (Donation/Volunteer/Complaint/etc.)
NERAgent	Extracts named entities using hybrid spaCy + Gemini approach
ResponseAgent	Generates a contextual professional email draft
RouterAgent	Maps intent to department and assigns priority level 1–4

The agents execute sequentially — Classify → Extract → Generate → Route — so that each agent has full context from all previous steps.

Layer 4 — Data Layer The data layer consists entirely of Firebase managed services. Firebase Firestore stores all triage results and user profiles as JSON documents in two root collections — triage_results and users. Firebase Authentication manages user identity, password hashing, and session tokens. The backend communicates with this layer exclusively through the Firebase Admin SDK — the frontend never has direct database access.

B. System Workflow- The complete end-to-end workflow of TriageAgent from user login to final response approval follows these steps:

Step 1 — Authentication The user visits the application at localhost:5173. The ProtectedRoute component checks for a valid JWT in localStorage. If no token is found, the user is redirected to the /login page. On successful login, Firebase Auth verifies the credentials, the backend issues a JWT access token, and the token is stored in localStorage via Zustand middleware. The user is redirected to the /home dashboard.

Step 2 — Message Submission The authenticated user navigates to the /triage page and submits a message through the MessageInput component — entering the message text (minimum 10 characters) and optionally providing sender name and sender email. On submission, the frontend sends a POST /api/triage request with the message payload and JWT Bearer token.

Step 3 — AI Triage Pipeline The FastAPI backend receives the request, validates the payload with Pydantic V2, verifies the JWT, and passes the validated MessageInput to the TriageOrchestrator. The Orchestrator generates a unique triage ID (TRG-{timestamp}-{uuid8}) and executes the four-agent pipeline:

- **ClassifierAgent** sends a structured prompt to Gemini 2.5 Flash (JSON mode, temperature=0.3) and receives urgency, intent, confidence, and reasoning.
- **NERAgent** runs spaCy on the message text (<50ms, offline) for standard entities, then calls Gemini LLM for domain-specific entities, and merges both outputs.
- **ResponseAgent** sends a structured prompt to Gemini 2.5 Flash (temperature=0.7, max_tokens=2000) and receives a complete draft email.
- **RouterAgent** performs a dictionary lookup (no LLM call) to map intent to department and urgency to priority level.

Step 4 — Result Storage & Return The Orchestrator assembles the complete TriageResult object, saves it to Firebase Firestore (triage_results/{triage_id}), and returns a HTTP 200 response with the full TriageResult JSON to the frontend. Total processing time: 5–10 seconds.

Step 5 — Human Review The frontend displays the triage result in a 3-tab interface. The staff member reviews the Overview tab (original message, AI reasoning, routing), the Entities tab (all extracted data points as color-coded pills), and the Draft Response tab (subject line + email body with copy button). The staff member clicks **Approve** or **Reject**.

Step 6 — Feedback & Status Update On approval or rejection, the frontend sends a POST /api/triage/{triage_id}/feedback request. The backend updates the document status in Firestore (approved or rejected) and timestamps the update. The result is removed from the pending review queue and reflected in the dashboard and analytics.

Step 7 — Dashboard & Analytics The staff member can view all processed messages on the /dashboard page with urgency and status filters, or visit /analytics to review performance KPIs, urgency distribution, and intent trends across configurable time periods.

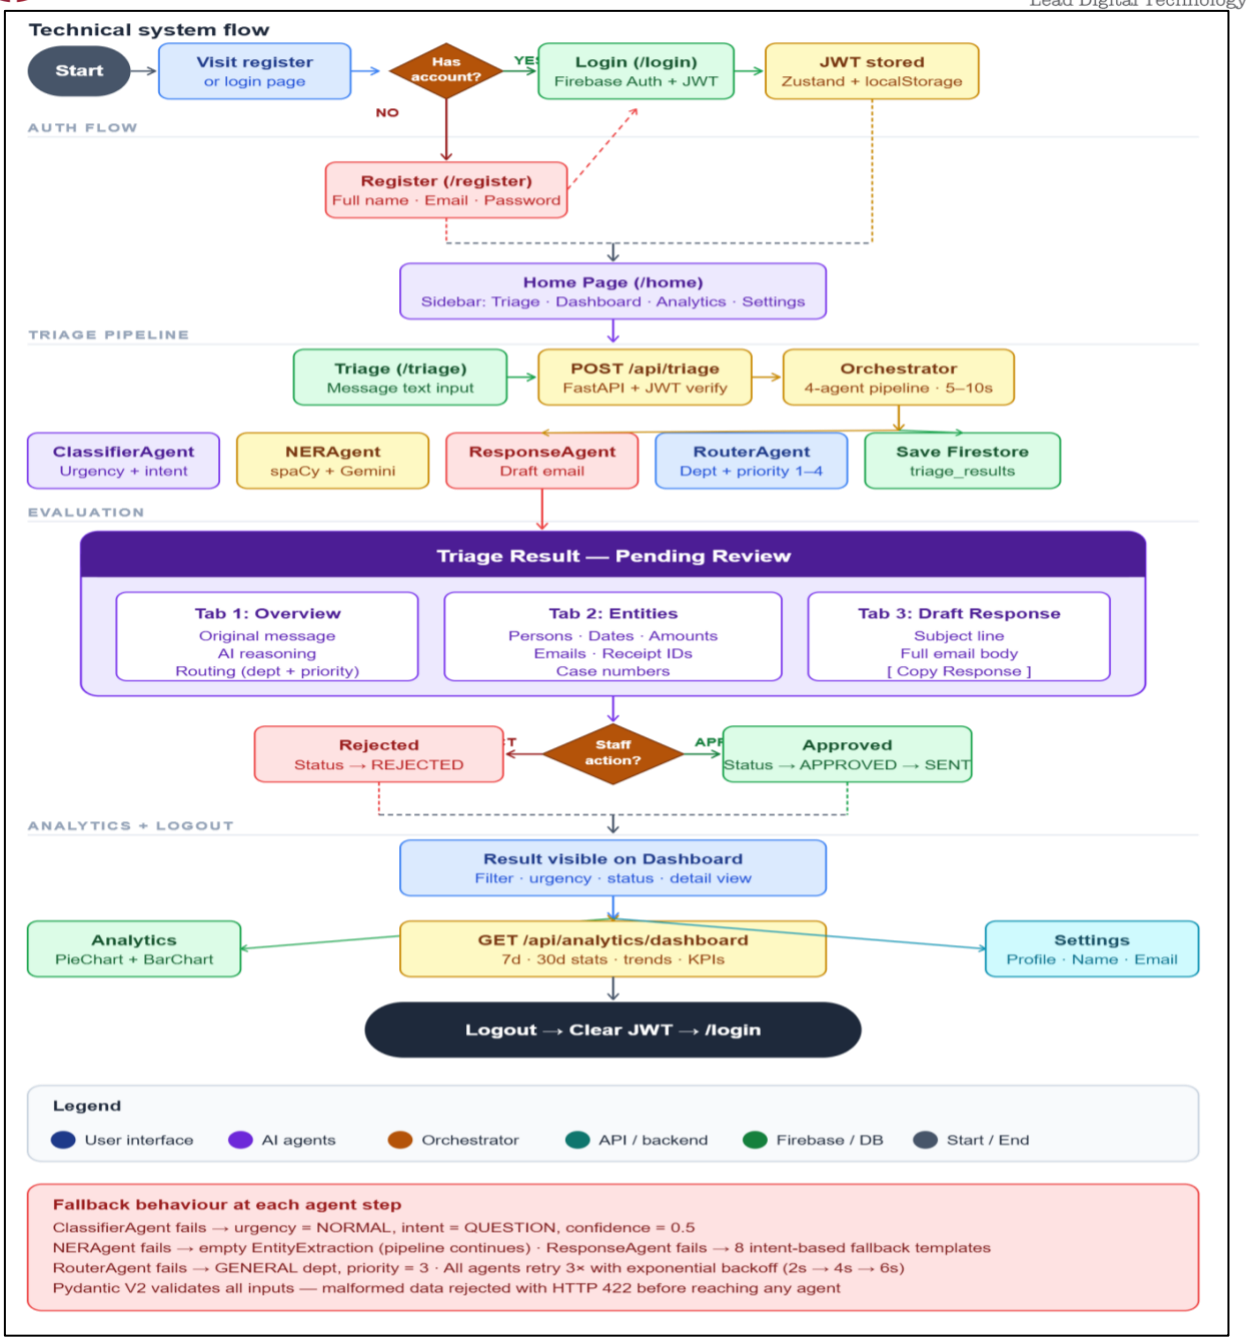

3. Tools & Technologies Used

TriageAgent is built using a modern, carefully selected technology stack that covers every layer of the application — from the user interface to the AI agent pipeline to cloud-based data storage. Each tool was chosen based on its specific strengths, compatibility with the project requirements, and suitability for building a scalable, secure, and AI-powered communication triage platform.

The frontend uses React-based technologies for a responsive and interactive staff dashboard, the backend uses Python with FastAPI for high-performance async API handling, and the AI layer integrates Gemini 2.5 Flash along with spaCy for intelligent message classification, entity extraction, and professional draft response generation.

Category	Technology	Version	Purpose
Frontend Framework	React 18 + Vite	React 19.2 / Vite 7.3.1	SPA framework with fast HMR and production bundling
UI Styling	Tailwind CSS 3	3.4.1	Utility-first CSS for all components and layouts
State Management	Zustand	5.0.11	Global state with localStorage JWT persistence
HTTP Client	Axios	1.13.6	REST API calls with JWT interceptor and 401 redirect
Animations	Framer Motion	12.38.0	Page transitions and component animations
Charts	Recharts	3.7.0	PieChart (urgency), BarChart (intent) on analytics page
Routing	React Router v7	7.13.1	Client-side routing with protected route enforcement
Backend Framework	FastAPI	0.109.0	Async REST API with auto-generated Swagger docs
Runtime	Python	3.11.0	Backend language for all agents, services, and API logic
ASGI Server	Uvicorn	0.27.0	High-performance async server for FastAPI
Agent Framework	LangChain	0.2.17	Multi-agent orchestration, prompt templates, output parsers
LLM API	Google Gemini 2.5 Flash	google-genai 1.46	Classification, entity extraction, response generation
NER (Offline)	spaCy	3.7.2	Fast offline entity extraction — names, dates, amounts
Data Validation	Pydantic V2	2.5.3	Request/response schema validation at every API boundary
Database	Firebase Firestore	firebase-admin 7.2	Cloud NoSQL document store for triage results and users

Authentication	Firebase Auth	firebase-admin 7.2	User identity, password hashing, credential management
API Authorization	JWT (python-jose)	3.3.0	Stateless API access token — HS256, 30-minute expiry
Version Control	Git + GitHub	Repository Link	Source code management and repository hosting
IDE	Visual Studio Code	—	Development environment with Python venv and npm

Together, these technologies form a cohesive and well-integrated stack that enables TriageAgent to process incoming non-profit communications intelligently, accurately, and securely. The combination of a stateless FastAPI backend, a dual-token authentication model, a hybrid NER approach, and a cloud-first Firebase infrastructure ensures the system is reliable, maintainable, and ready for future scalability into a full production deployment on Google Cloud Platform.

Results & Output

1. Screenshots / outputs

Screen 1: Frontend Development Server (Vite) - The TriageAgent frontend development server running via Vite v7.3.1 on localhost:5173. The terminal confirms the React application is successfully compiled and running.

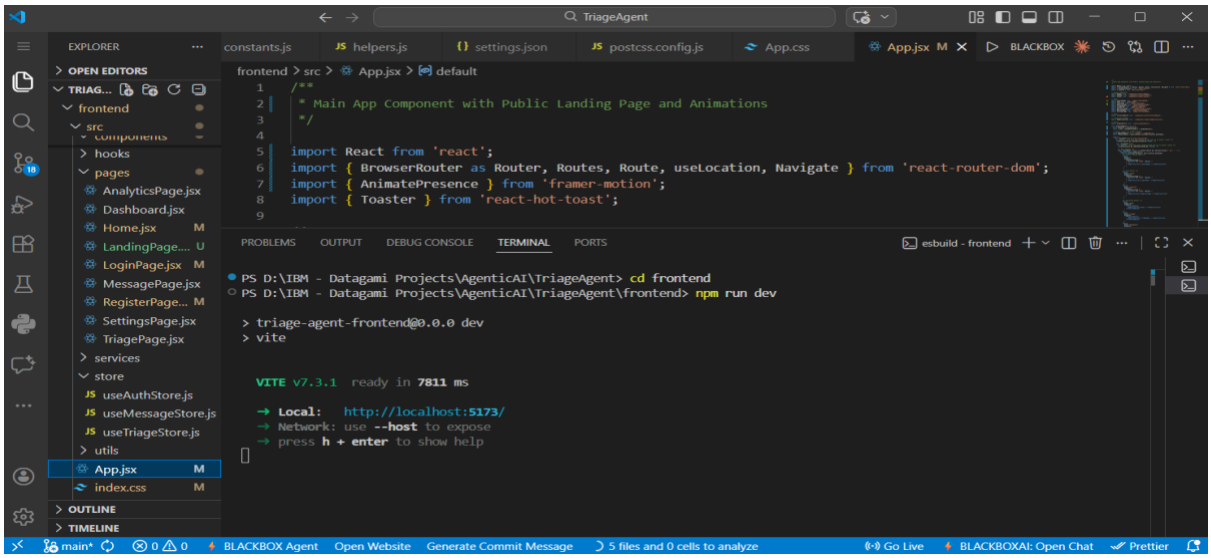

Screen 2: Backend Server & API Logs (Uvicorn + FastAPI) - The FastAPI backend server running via Uvicorn on http://0.0.0.0:8000 with Python 3.11. The terminal confirms successful initialization — including the Gemini AI services, Agents, Spacy, Firebase Services, and shows live API request logs with HTTP 200 OK responses for endpoints.

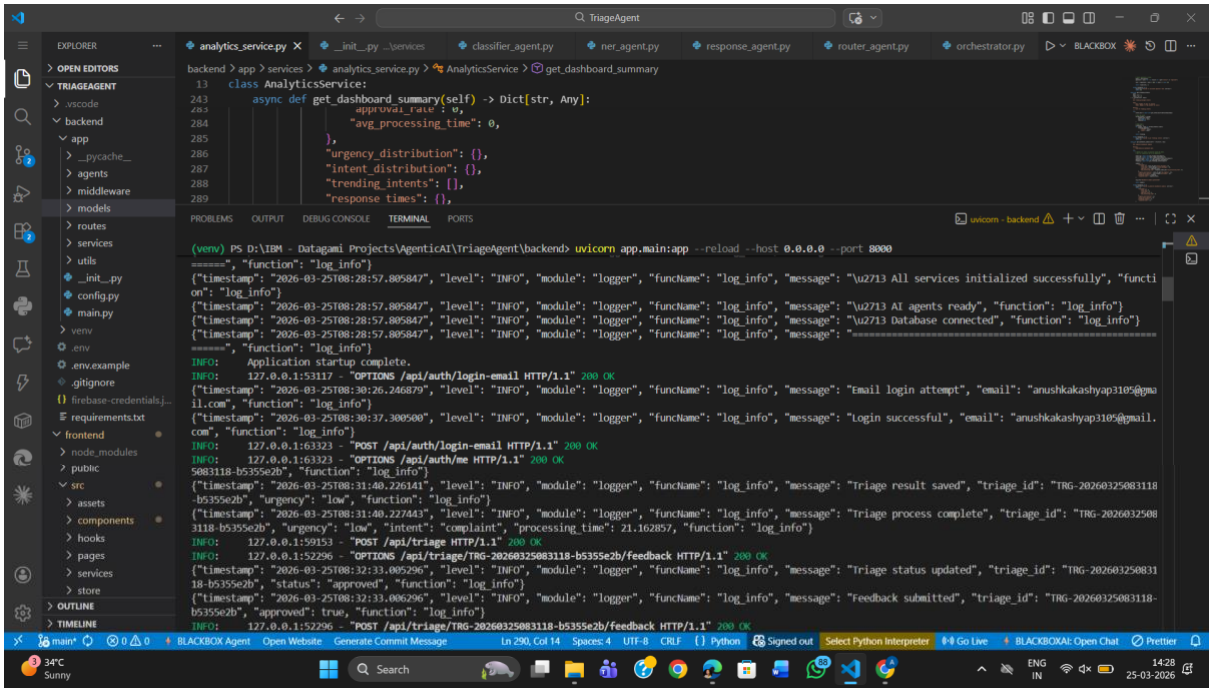

Screen 3: Firebase Firestore Console - The Firestore console showing the triage_results collection with real triage documents containing embedded classification, entities, routing MAPs, draft response text, status, and timestamps — confirming successful database persistence.

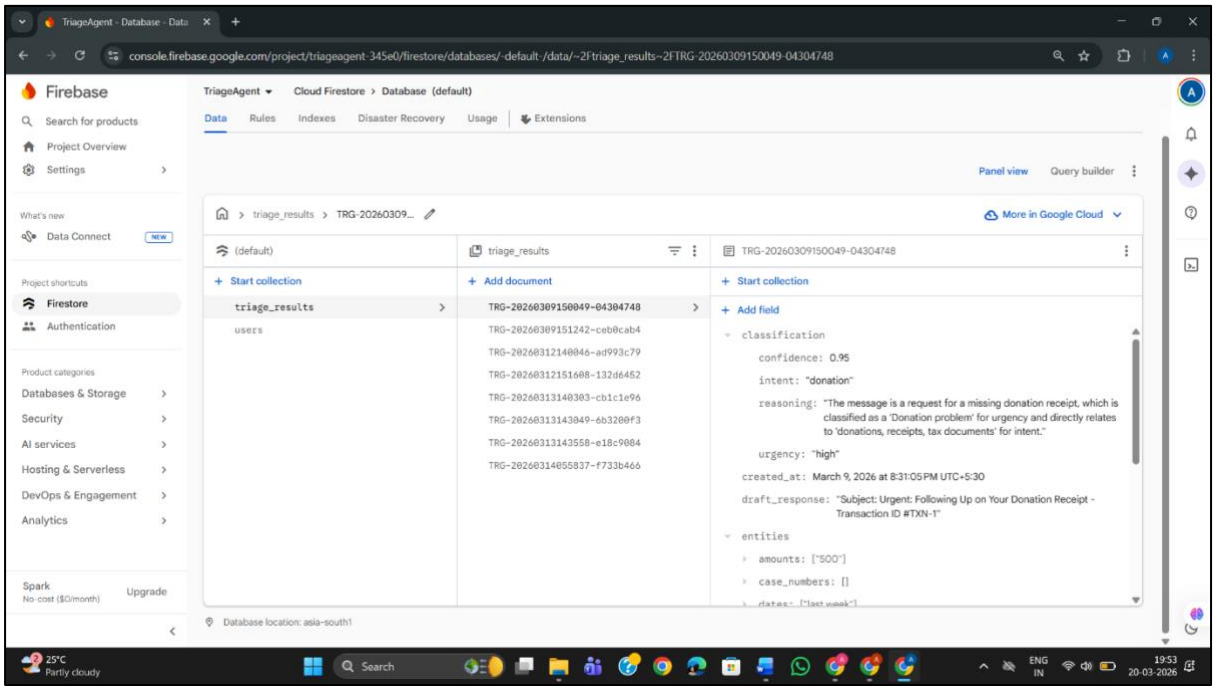

Screen 4: Landing Page - The TriageAgent public landing page displaying the animated hero section with logo, headline, CTA buttons, Introduction and how it works, a stats section, and a 6-card features grid, and all animated on scroll using Framer Motion.

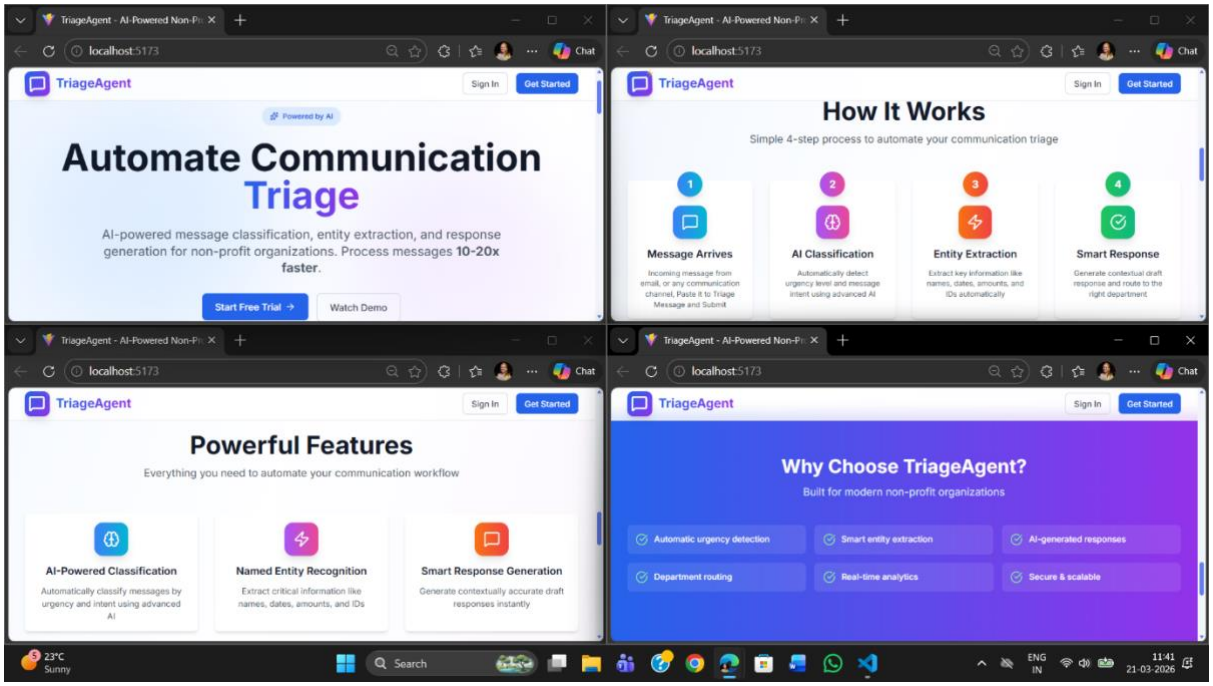

Screen 5: User Registration Page - The registration interface where new users can create an account using their name, email address, and password via Firebase Authentication.

Screen 6: User Login Page - The login page where registered users authenticate using their email and password. Firebase verifies credentials and the backend issues a JWT access token stored in localStorage.

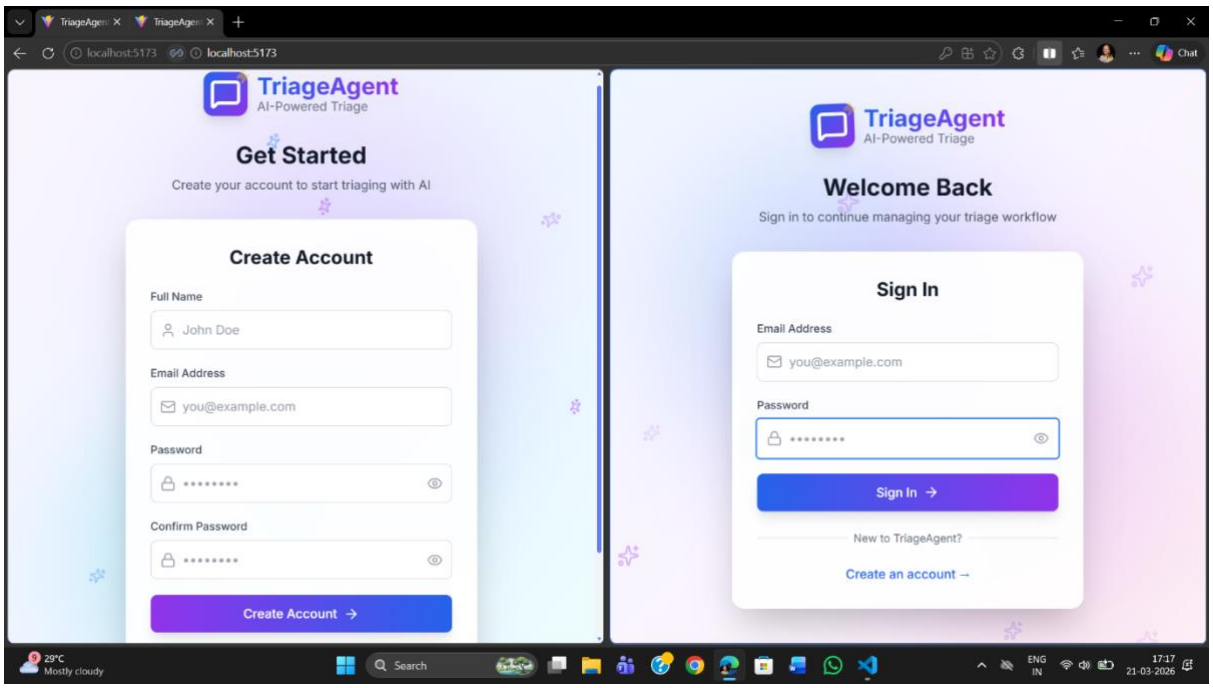

Screen 7: Home Page (Post-Login) - After authentication, users land on the Home page featuring the full sidebar navigation, header with user email and logout button, feature overview cards, and key system statistics.

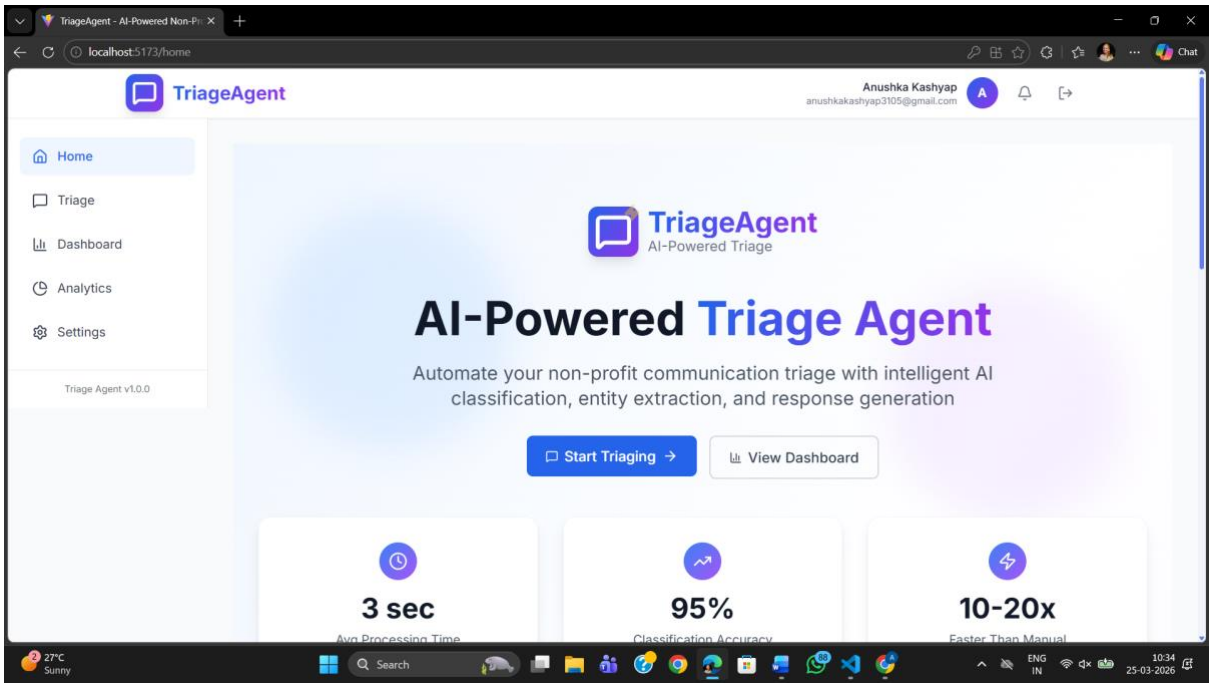

Screen 8: Triage Page — Input Form - The triage interface displays a clean message input form with a text area (min 10 characters), optional sender email and sender name fields, a Triage Message submit button with loading state, and a Clear button. A How It Works info banner explains the 5-step process below the form.

Screen 9: Triage Page — Result (Tab 1: Overview) - After processing, the result displays the triage ID, status badge (Pending Review), urgency badge, intent badge, confidence percentage, AI reasoning explanation, routing information, and the original message text.

Screen 10: Triage Page — Result (Tab 2: Entities) - The entities tab shows all extracted data points displayed as color-coded pill badges grouped by type — People, Organizations, Dates, Amounts, Emails, Phone Numbers, and Receipt IDs.

Screen 11: Triage Page — Result (Tab 3: Draft Response) - The draft response tab shows the AI-generated subject line followed by the complete email body, and a Copy Response button that copies the full response to clipboard with a toast notification.

Screen 12: Approve / Reject Actions - Below the result tabs, two action buttons — Approve Response (green) and Reject & Edit (red) — allow staff user to submit feedback. and update the triage status in Firestore.

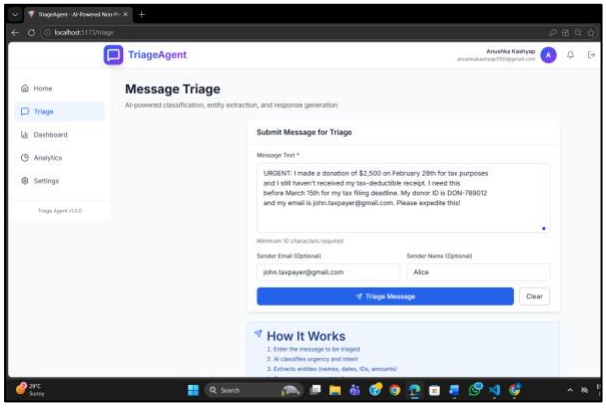

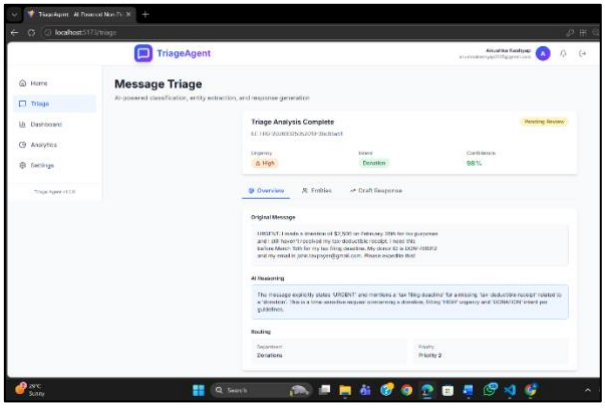

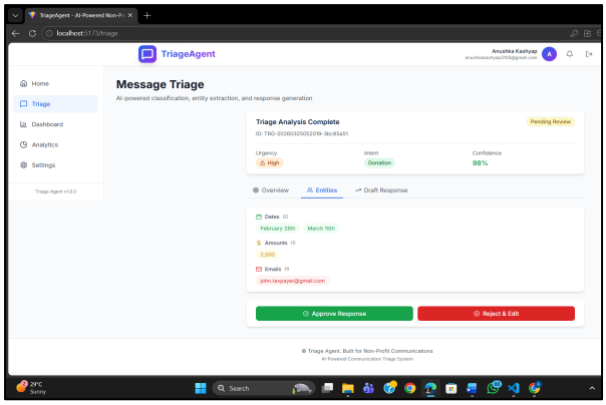

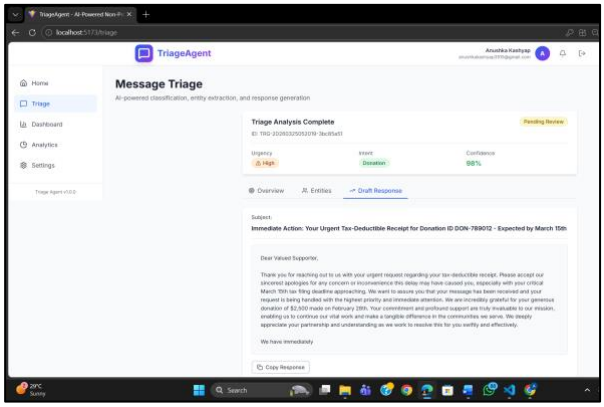

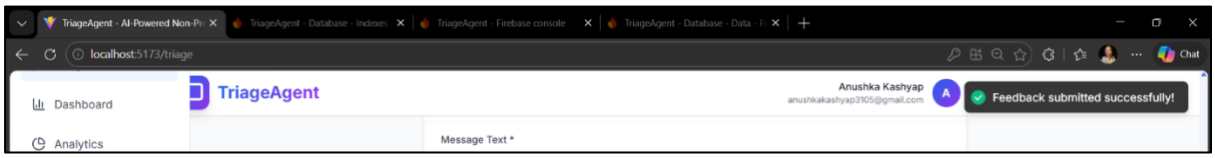

Screen 13: Dashboard Page - The dashboard displays a filterable, paginated list of all triaged messages where each card shows urgency badge, intent badge, status badge, message preview, sender info, timestamp, and assigned department.

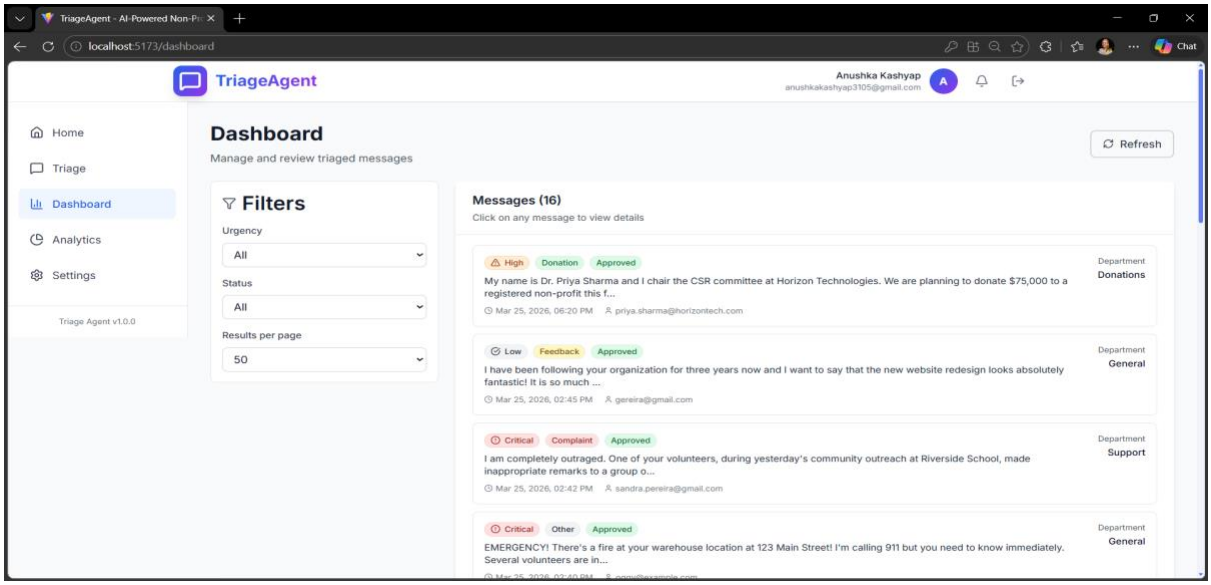

Screen 14: Analytics Page - The analytics page displays four KPI StatsCards, a PieChart showing urgency distribution, and a BarChart showing intent breakdown.

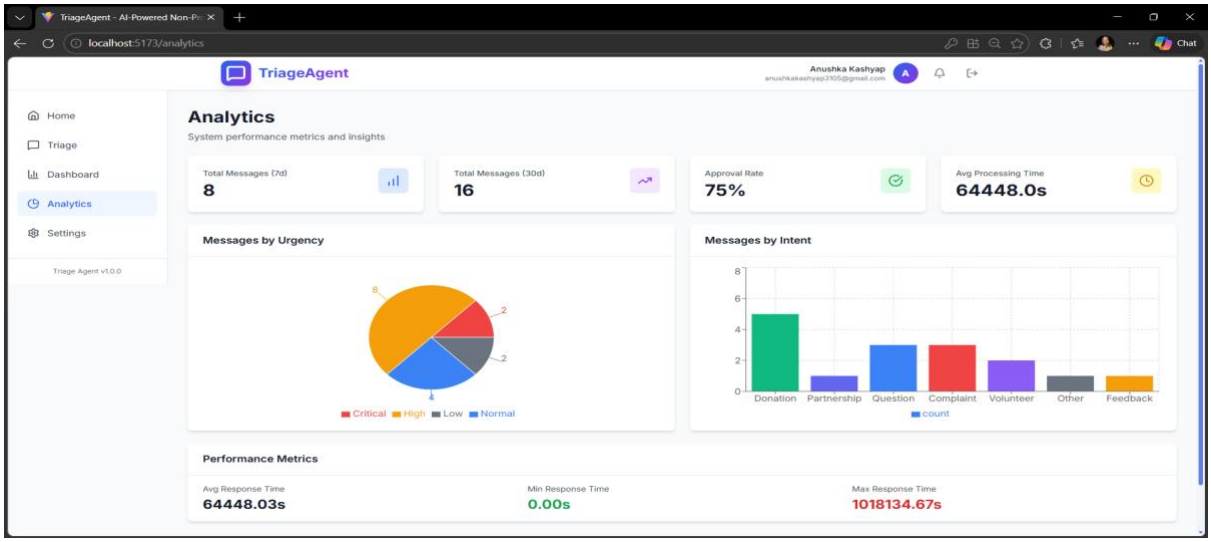

Screen 15: SettingsPage - The settings page shows the authenticated user's profile card.

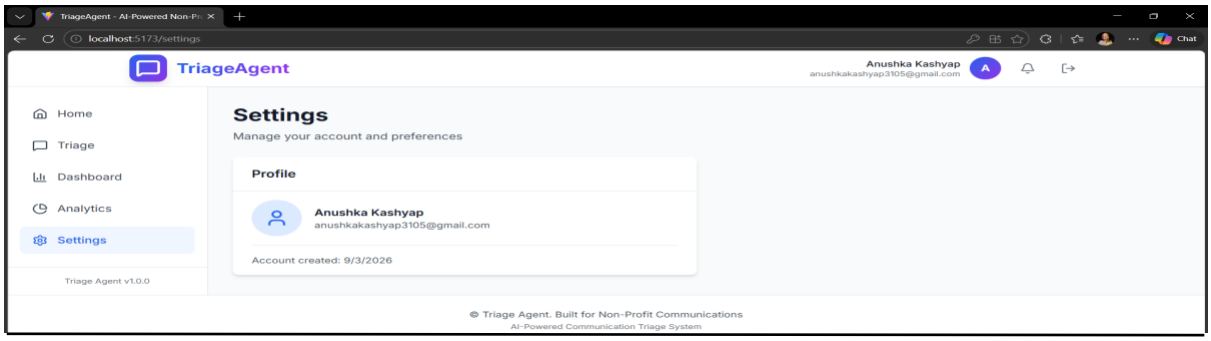

2. Reports / Dashboards / Models

A. Reports

- **Triage Result Report** Every processed message generates a complete structured TriageResult document stored in Firestore with the following fields — triage_id, original message text, classification output (urgency, intent, confidence, reasoning), all extracted entities (8 entity type arrays), the AI-generated draft_response, routing (department, priority, suggested_assignee), review status, created_at, updated_at, and metadata (sender info, processing time). This serves as a permanent audit trail for every communication the organization received.
- **API Documentation Report** The FastAPI Swagger UI at localhost:8000/docs serves as a live, auto-generated API reference report. It documents all 13 endpoints with full request body schemas, response models, authentication requirements, query parameter descriptions, and interactive test execution — making it the single source of truth for all API contracts in the system.

B. Dashboard

- **Admin Triage Dashboard** The primary operational dashboard gives non-profit staff a complete view of all incoming communications in the system. It provides real-time filtering by urgency (Critical / High / Normal / Low) and status (Pending Review / Approved / Rejected / Sent / Archived), paginated message list with per-row urgency and intent indicators, one-click navigation to individual message detail views, and a Refresh button for fetching the latest data from Firestore.
- **Analytics Dashboard** The analytics dashboard serves as the management reporting layer of the system. It aggregates all triage data stored in Firestore and reports it through visual components — KPI cards for headline metrics, a PieChart for urgency distribution, a BarChart for intent breakdown, and performance metrics (average, minimum, and maximum processing times, plus approval rate percentage). The dashboard supports configurable time period filters (7 days, 30 days, 90 days).

C. Models: The system uses four AI agent models working in sequence:

- **AI Classification Model Output** It uses Gemini 2.5 Flash model. Every classification produced by the ClassifierAgent includes four fields — urgency (UrgencyLevel enum), intent (IntentCategory enum), confidence (float 0–1), and reasoning (LLM-generated explanation). Confidence scores above 0.85 indicate high certainty. The model uses temperature=0.3 for consistency and JSON mode to ensure structured, parseable output every time.

- **Entity Extraction Model Output** It uses spaCy + Gemini 2.5 Flash hybrid model. The NERAgent produces an EntityExtraction object with 8 typed arrays — persons, organizations, dates, amounts, emails, phone_numbers, receipt_ids, case_numbers. All arrays are deduplicated using Python set(). The hybrid approach ensures both standard NLP entities and domain-specific non-profit identifiers are captured.
- **Draft Response Model Output** It uses Gemini 2.5 Flash model. The ResponseAgent produces a plain text email draft with a Subject: line, greeting, body paragraphs addressing the specific issue, and a closing signature. The tone is automatically adjusted based on urgency — Critical messages use urgent empathetic language, Normal messages use standard professional tone. Drafts average 150–300 words. The model uses temperature=0.7, max_tokens=2000 with structured prompt.
- **Intelligent Routing Model Output** The RouterAgent uses a pure dictionary lookup model with no LLM call. It maps IntentCategory → DepartmentType using INTENT_TO_ DEPARTMENT and UrgencyLevel → priority integer using URGENCY_PRIORITY — both defined in enums.py. Since routing is fully deterministic, it is instantaneous (<1ms) and consistent regardless of API availability. Output is a RoutingDecision object containing department, priority level, and suggested assignee role.

3. Key Outcomes

- Successfully built a fully functional AI-powered communication triage system for non-profit organizations, covering the complete workflow from message input to human-reviewed draft response.
- Dual-token authentication (Firebase Auth + JWT) provides secure, stateless session management with automatic token attachment, 401 auto-logout, and localStorage persistence.
- Google Gemini 2.5 Flash accurately classifies incoming messages by urgency and intent with 90–95% accuracy, providing confidence scores and AI reasoning for every result.
- Hybrid NER (spaCy + Gemini LLM) successfully extracts both standard entities (names, dates, amounts) and domain-specific identifiers (receipt IDs, numbers).
- AI-generated draft email responses with subject lines and complete email bodies are produced in 2–3s, tailored to message content, urgency, and extracted entities
- Intelligent department routing with priority assignment (1–4) achieved through deterministic dictionary lookup — instantaneous, consistent, and API-independent.

- Complete AI triage pipeline completes in 5–10 seconds per message — a 10–20× improvement over the 15–20 minute manual triage process.
- Human-in-the-loop review workflow fully operational — every AI response requires explicit staff approval before being marked as sent.
- All triage results reliably persisted in Firebase Firestore as structured documents, serving as a permanent audit trail for every communication received.
- Exponential backoff retry logic (3 attempts: 2s → 4s → 6s) ensures system reliability under Gemini API errors, with fallback responses guaranteeing the pipeline never crashes.
- Stateless FastAPI backend and cloud-hosted Firebase Firestore allow horizontal scaling with zero infrastructure changes.
- Modular 4-agent pipeline (Classifier → NER → Response → Router) — each agent independently testable, replaceable, and upgradeable without affecting the rest.
- Production-quality frontend with Framer Motion animations, animated landing page, color-coded urgency badges, tabbed result interface, and responsive layout.

Conclusion

TriageAgent successfully demonstrates how modern AI technologies can be combined to solve a real operational problem faced by non-profit organizations. It demonstrates AI-powered automation can meaningfully transform how non-profit organizations manage their incoming communications. By combining a multi-agent AI architecture with a modern full-stack implementation, the system delivers a complete, production-grade triage workflow that reduces manual processing time by 10–20× while maintaining full human oversight over every outgoing response.

The project fulfills all core requirements of the problem statement — real-time message classification, hybrid named entity recognition, contextual draft response generation, and intelligent department routing — within a secure, scalable, and reliable system. The human-in-the-loop design ensures that AI augments staff capability rather than replacing human judgment, which is especially critical in the empathy-driven context of non-profit communication.

The carefully designed 4-agent pipeline (Classifier → NER → Response → Router), coordinated by a central Orchestrator and backed by Google Gemini 2.5 Flash, proves that complex communication workflows requiring 15–20 minutes of manual effort can be reliably automated in 5–10 seconds. The hybrid NER approach combining spaCy offline processing with Gemini LLM contextual extraction demonstrates how different AI techniques can be combined to achieve results neither could deliver independently.

From a technical perspective, the system successfully integrates across all architectural layers — React 18 frontend, FastAPI backend, Google Gemini 2.5 Flash as the core LLM, spaCy for offline NER, LangChain for multi-agent orchestration, and Firebase for cloud infrastructure — into a cohesive, working application with 13 REST endpoints, real-time analytics, and a complete audit trail of every processed message. The fallback-first design, exponential backoff retry logic, Pydantic V2 input validation, and JWT-based authentication collectively ensure the system meets production-grade reliability and security standards.

The React 18 frontend with Framer Motion animations, real-time analytics dashboard, and intuitive 3-tab triage result interface delivers a professional user experience, making the system practical and immediately usable by non-profit staff without technical expertise.

In conclusion, TriageAgent is not just a functional proof of concept — it is a deployable, extensible, and maintainable system that solves a real operational problem for non-profit organizations.

Key learnings from this project include:

- Hands-on experience designing and implementing a multi-agent AI architecture using LangChain, with four specialized agents coordinated by a central Orchestrator.
- Practical integration of Google Gemini 2.5 Flash for message classification, named entity extraction, and contextual draft response generation in a real-world application.
- Implementing a hybrid NER pipeline combining spaCy offline processing with Gemini LLM — understanding when to use rule-based vs. AI-based approaches for optimal performance.
- Building production-grade LLM reliability through exponential backoff retry logic, fallback response templates, and defensive prompt engineering with JSON mode.
- Practical implementation of a dual-token authentication model using Firebase Authentication and JWT for secure, stateless API access control.
- Building a secure REST API with FastAPI including JWT middleware, CORS policy, rate limiting, Pydantic V2 input validation, and structured JSON logging.
- Working with Firebase Firestore as a cloud NoSQL database — designing document structures, performing filtered queries, and managing a complete status lifecycle.
- Designing and building a responsive Single Page Application with React 18, Zustand state management, Tailwind CSS, and Framer Motion animations.
- Understanding the importance of human-in-the-loop design — AI accelerates the process but human judgment must remain in control of every final decision.

Overall, TriageAgent successfully achieves its goal of creating an intelligent, production-grade communication triage system that bridges the gap between high-volume incoming messages and efficient, empathetic non-profit staff responses.

Future Scope & Enhancements

As a v1.0 system running in a local development environment, TriageAgent has a strong foundation for future enhancements. The following improvements are planned for future versions:

- **Direct Channel Integrations** Integrate natively with email providers (Gmail, Outlook), WhatsApp Business API, SMS (Twilio), and web contact forms — allowing messages to flow directly into the triage pipeline without manual copy-paste.
- **Multi-Language Support** Extend the classification, NER, and response generation pipeline to support non-English languages — enabling non-profit organizations operating in multilingual regions to use the system effectively.
- **Role-Based Access Control (RBAC)** Introduce multiple user roles — Admin, Manager, Staff, Viewer — with different permission levels for reviewing messages, approving responses, accessing analytics, and managing system settings.
- **Cloud Deployment** Deploy the FastAPI backend to Google Cloud Run for auto-scaling containerized hosting, and the React frontend to Firebase Hosting — making the system publicly accessible with zero infrastructure management.
- **Real-Time Notifications** Implement WebSocket or Firebase real-time listeners to push live notifications to staff when new Critical or High urgency messages arrive — ensuring no urgent communication goes unnoticed.
- **Advanced Analytics & Reporting** Add exportable PDF/CSV reports, trend forecasting, response time SLA tracking, staff performance metrics, and custom date range filters to the analytics dashboard.
- **Redis Caching Layer** Introduce Redis for caching frequently accessed analytics queries and message list results — reducing Firestore read operations and improving dashboard load times at scale.
- **Expanded Agent Capabilities** Add a Sentiment Analysis Agent to detect emotional tone, a Duplicate Detection Agent to flag repeated messages from the same sender, and a Priority Escalation Agent to automatically escalate unreviewed Critical messages after a defined time window.